

Problem Set #2

Anaya Hall & Christian Miller

Part 1: Theory

Part 2: Applied: Returns to Scale in Electricity Supply

First, load our OLS function created in Problem Set #1. We're including a built in t-test this time around.

```
ols <- function(data, y_data, X_data, intercept = T, H0 = 0, two_tail = T, alpha = 0.05) {  
  # Function setup ----  
  # Require the 'dplyr' package  
  require(dplyr)  
  # Function to convert tibble, data.frame, or tbl_df to matrix  
  to_matrix <- function(the_df, vars) {  
    # Create a matrix from variables in var  
    new_mat <- the_df %>%  
      # Select the columns given in 'vars'  
      select_(.dots = vars) %>%  
      # Convert to matrix  
      as.matrix()  
    # Return 'new_mat'  
    return(new_mat)  
  }  
  
  # Create dependent and independent variable matrices ----  
  # y matrix  
  y <- to_matrix (the_df = data, vars = y_data)  
  # X matrix  
  X <- to_matrix (the_df = data, vars = X_data)  
  # If 'intercept' is TRUE, then add a column of ones  
  if (intercept == T) {  
    X <- cbind(1,X)  
    colnames(X) <- c("intercept", X_data)  
  }  
  
  # Calculate b, y_hat, and residuals ----  
  b <- solve(t(X) %*% X) %*% t(X) %*% y  
  y_hat <- X %*% b  
  e <- y - y_hat  
  
  # Useful -----  
  n <- nrow(X) # number of observations  
  k <- ncol(X) # number of independent variables  
  dof <- n - k # degrees of freedom  
  i <- rep(1,n) # column of ones for demeaning matrix  
  A <- diag(i) - (1 / n) * i %*% t(i) # demeaning matrix  
  y_star <- A %*% y # for SST  
  X_star <- A %*% X # for SSM  
  SST <- drop(t(y_star) %*% y_star)
```

```

SSM <- drop(t(b) %*% t(X_star) %*% X_star %*% b)
SSR <- drop(t(e) %*% e)

# Measures of fit and estimated variance ----
R2uc <- drop((t(y_hat) %*% y_hat)/(t(y) %*% y)) # Uncentered R^2
R2 <- 1 - SSR/SST # Uncentered R^2
R2adj <- 1 - (n-1)/dof * (1 - R2) # Adjusted R^2
AIC <- log(SSR/n) + 2*k/n # AIC
SIC <- log(SSR/n) + k/n*log(n) # SIC
s2 <- SSR/dof # s^2

# Measures of fit table ----
mof_table_df <- data.frame(R2uc, R2, R2adj, SIC, AIC, SSR, s2)
mof_table_col_names <- c("$R^2_\\text{uc}$", "$R^2$",
                        "$R^2_\\text{adj}$",
                        "SIC", "AIC", "SSR", "$s^2$")
mof_table <- mof_table_df %>% knitr::kable(
  row.names = F,
  col.names = mof_table_col_names,
  format.args = list(scientific = F, digits = 4),
  booktabs = T,
  escape = F
)

# t-test----
# Standard error
se <- as.vector(sqrt(s2 * diag(solve(t(X) %*% X))))
# Vector of _t_ statistics
t_stats <- (b - H0) / se
# Calculate the p-values
if (two_tail == T) {
  p_values <- pt(q = abs(t_stats), df = dof, lower.tail = F) * 2
} else {
  p_values <- pt(q = abs(t_stats), df = dof, lower.tail = F)
}
# Do we (fail to) reject?
reject <- ifelse(p_values < alpha, reject <- "Reject", reject <- "Fail to Reject")

# Nice table (data.frame) of results
ttest_df <- data.frame(
  # The rows have the coef. names
  effect = rownames(b),
  # Estimated coefficients
  coef = as.vector(b) %>% round(3),
  # Standard errors
  std_error = as.vector(se) %>% round(3),
  # t statistics
  t_stat = as.vector(t_stats) %>% round(3),
  # p-values
  p_value = as.vector(p_values) %>% round(4),
  # reject null?
  significance = as.character(reject)
)

```

```

ttest_table <- ttest_df %>% knitr::kable(
  booktabs = T,
  format.args = list(scientific = F),
  escape = F
)

# Data frame for exporting for y, y_hat, X, and e vectors ----
export_df <- data.frame(y, y_hat, e, X) %>% tbl_df()
colnames(export_df) <- c("y", "y_hat", "e", colnames(X))

# Return ----
return(list(n=n, dof=dof, b=b, vars=export_df, R2uc=R2uc, R2=R2,
  R2adj=R2adj, AIC=AIC, SIC=SIC, s2=s2, SST=SST, SSR=SSR,
  mof_table=mof_table, ttest=ttest_table))
}

```

We'll also need a function to do an F-test for this Problem Set.

```

# Joint test function (from Ed's notes). Could also write a more complex functions that takes R and r.
F_test <- function(data, y_var, X_vars) {
  # Turn data into matrices
  y <- to_matrix(data, y_var)
  X <- to_matrix(data, X_vars)
  # Add intercept
  X <- cbind(1, X)
  # Name the new column "intercept"
  colnames(X) <- c("intercept", X_vars)
  # Calculate n and k for degrees of freedom
  n <- nrow(X)
  k <- ncol(X)
  # J is k-1
  J <- k - 1
  # Create the R matrix: bind a column of zeros
  # onto a J-by-J identity matrix
  R <- cbind(0, diag(J))

  # Estimate coefficients
  b <- b_ols(data, y_var, X_vars)
  # Calculate OLS residuals
  e <- y - X %*% b
  # Calculate s^2
  s2 <- (t(e) %*% e) / (n-k)
  # Force s2 to numeric
  s2 %<>% as.numeric()

  # Create the inner matrix  $R(X'X)^{-1}R'$ 
  RXXR <- R %*% solve(t(X) %*% X) %*% t(R)
  # Calculate the F stat
  f_stat <- t(R %*% b) %*% solve(RXXR) %*% (R %*% b) / J / s2
  # Calculate the p-value
  p_value <- pf(q = f_stat, df1 = J, df2 = n-k, lower.tail = F)
  # Create a data.frame of the f stat. and p-value
  results <- data.frame(
    f_stat = f_stat %>% as.vector(),

```

```

    p_value = p_value %>% as.vector()
  return(results)
}

```

Question 1:

Read the data into R. Inspect it. Sort by size (Q (kwh)).

```

nerlove <- readxl::read_excel("nerlove.xls", col_names=T)

# Fix typo in 13th row (missing a decimal!)
# DO THIS MORE ELEGANTLY!
nerlove[13, "PL"] <- 1.81

nerlove %>% arrange(Q)

```

```

## # A tibble: 145 x 5
##       TC      Q    PL    PF    PK
##   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.0820     2.  2.09  17.9  183.
## 2 0.661      3.  2.05  35.1  174.
## 3 0.990      4.  2.05  35.1  171.
## 4 0.315      4.  1.83  32.2  166.
## 5 0.197      5.  2.12  28.6  233.
## 6 0.0980     9.  2.12  28.6  195.
## 7 0.949     11.  1.98  35.5  206.
## 8 0.675     13.  2.05  35.1  150.
## 9 0.525     13.  2.19  29.1  155.
## 10 0.501     22.  1.72  15.0  188.
## # ... with 135 more rows

```

```
head(nerlove)
```

```

## # A tibble: 6 x 5
##       TC      Q    PL    PF    PK
##   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.0820     2.  2.09  17.9  183.
## 2 0.661      3.  2.05  35.1  174.
## 3 0.990      4.  2.05  35.1  171.
## 4 0.315      4.  1.83  32.2  166.
## 5 0.197      5.  2.12  28.6  233.
## 6 0.0980     9.  2.12  28.6  195.

```

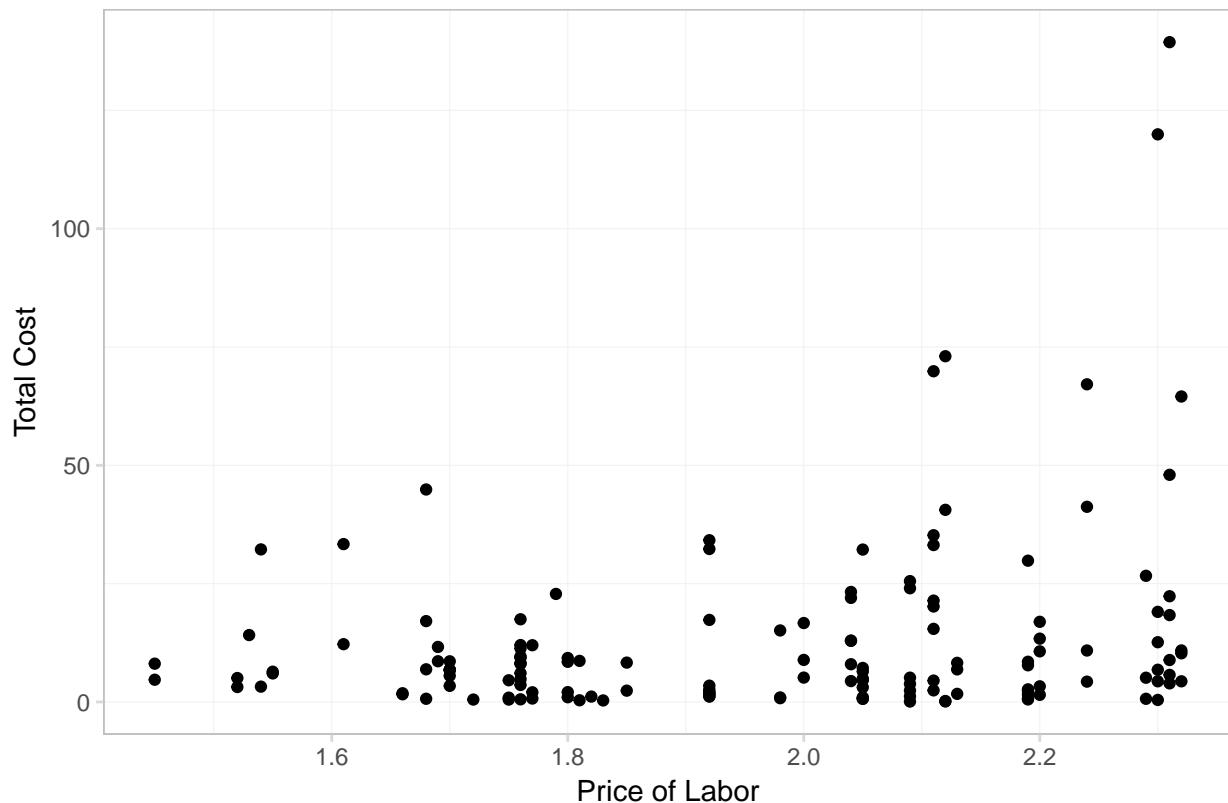
Plot the series and make sure your data are read in correctly.

```

ggplot(nerlove, aes(x=PL, y=TC)) +
  geom_point() +
  labs(title="Nerlove Data", x="Price of Labor", y="Total Cost") +
  theme_are

```

Nerlove Data



```
### plot rest of independent vars too.
```

Question 2:

Replicate regression I (page 176) in the paper.

Regression I:

$$\log(TC) - \log(P_F) = \beta_0 + \beta_1 Q + \beta_2 (\log(P_L) - \log(P_F)) + \beta_3 (\log(P_K) - \log(P_F))$$

Equivalent to:

$$\log\left(\frac{TC}{P_F}\right) = \beta_0 + \beta_1 Q + \beta_2 \left(\log\left(\frac{P_L}{P_F}\right)\right) + \beta_3 \left(\log\left(\frac{P_K}{P_F}\right)\right)$$

Where:

TC = total production cost, P_L = wage rate, P_K = “price” of capital, P_F = price of fuel

Q = output (measured in kWh)

In generalized Cobb-Douglas form: $\beta_1 = \frac{1}{r}$, $\beta_2 = \frac{a_L}{r}$, $\beta_3 = \frac{a_K}{r}$

Prepare variables for Regression I.

```
# Create log variables
nerlove %<>% mutate(
  TClog = log10(TC),
  Qlog = log10(Q),
  PLlog = log10(PL),
  PKlog = log10(PK),
  PFlog = log10(PF)
```

```
)

# Create PF scaled variables
nerlove %<>% mutate(
  TCscaled = TClog - PFlog,
  PLscaled = PLlog - PFlog,
  PKscaled = PKlog - PFlog
)
```

Variable names: $\log(\frac{TC}{P_F})$ = "TCscaled" $\log(\frac{P_L}{P_F})$ = "PLscaled" $\log(\frac{P_K}{P_F})$ = "PKscaled"

```
# Regression I:
# dep var = (log costs - log fuel price) = TCscaled
reg_I <- ols(data = nerlove, y_data = "TCscaled",
             X_data = c("Qlog", "PLscaled", "PKscaled"),
             intercept = T, H0 = 0, alpha = 0.05)

reg_I$tttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-2.037	0.384	-5.301	0.0000	Reject
Qlog	0.721	0.017	41.334	0.0000	Reject
PLscaled	0.593	0.205	2.898	0.0044	Reject
PKscaled	-0.007	0.191	-0.039	0.9692	Fail to Reject

```
reg_I$mof_table
```

	R^2_{uc}	R^2	R^2_{adj}	SIC	AIC	SSR	s^2
	0.966	0.9316	0.9301	-3.433	-3.515	4.082	0.02895
Coefficients are pretty close to those in the paper. R^2 match!							

Question 3:

Conduct the hypothesis test using constant returns to scale ($\beta_1 = 1$) as your null hypothesis. What is the p -value associated with your test statistic? What is your point estimate of returns to scale? Constant? Increasing? Decreasing?

```
# Re-run regression with null hypothesis H0 = 1
reg_I <- ols(data = nerlove, y_data = "TCscaled",
             X_data = c("Qlog", "PLscaled", "PKscaled"),
             intercept = T, H0 = 1, alpha = 0.05)

reg_I$tttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-2.037	0.384	-7.903	0.0000	Reject
Qlog	0.721	0.017	-16.020	0.0000	Reject
PLscaled	0.593	0.205	-1.990	0.0485	Reject
PKscaled	-0.007	0.191	-5.282	0.0000	Reject