

Problem Set #2

Anaya Hall & Christian Miller

Part 1: Theory

(for practice only)

Part 2: Applied - Returns to Scale in Electricity Supply

First, load our OLS function created in Problem Set #1. We're including a built in t-test this time around.

```
ols <- function(data, y_data, X_data, intercept = T, H0 = 0, two_tail = T, alpha = 0.05) {  
  # Function setup ----  
  # Require the 'dplyr' package  
  require(dplyr)  
  # Function to convert tibble, data.frame, or tbl_df to matrix  
  to_matrix <- function(the_df, vars) {  
    # Create a matrix from variables in var  
    new_mat <- the_df %>%  
      # Select the columns given in 'vars'  
      select_(.dots = vars) %>%  
      # Convert to matrix  
      as.matrix()  
    # Return 'new_mat'  
    return(new_mat)  
  }  
  
  # Create dependent and independent variable matrices ----  
  # y matrix  
  y <- to_matrix(the_df = data, vars = y_data)  
  # X matrix  
  X <- to_matrix(the_df = data, vars = X_data)  
  # If 'intercept' is TRUE, then add a column of ones  
  if (intercept == T) {  
    X <- cbind(1, X)  
    colnames(X) <- c("intercept", X_data)  
  }  
  
  # Calculate b, y_hat, and residuals ----  
  b <- solve(t(X) %*% X) %*% t(X) %*% y  
  y_hat <- X %*% b  
  e <- y - y_hat  
  
  # Useful -----  
  n <- nrow(X) # number of observations  
  k <- ncol(X) # number of independent variables  
  dof <- n - k # degrees of freedom
```

```

i <- rep(1,n) # column of ones for demeaning matrix
A <- diag(i) - (1 / n) * i %*% t(i) # demeaning matrix
y_star <- A %*% y # for SST
X_star <- A %*% X # for SSM
SST <- drop(t(y_star) %*% y_star)
SSM <- drop(t(b) %*% t(X_star) %*% X_star %*% b)
SSR <- drop(t(e) %*% e)

# Measures of fit and estimated variance ----
R2uc <- drop((t(y_hat) %*% y_hat)/(t(y) %*% y)) # Uncentered R^2
R2 <- 1 - SSR/SST # Uncentered R^2
R2adj <- 1 - (n-1)/dof * (1 - R2) # Adjusted R^2
AIC <- log(SSR/n) + 2*k/n # AIC
SIC <- log(SSR/n) + k/n*log(n) # SIC
s2 <- SSR/dof # s^2

# Measures of fit table ----
mof_table_df <- data.frame(R2uc, R2, R2adj, SIC, AIC, SSR, s2)
mof_table_col_names <- c("$R^2_\\text{uc}$", "$R^2$",
                        "$R^2_\\text{adj}$",
                        "SIC", "AIC", "SSR", "$s^2$")
mof_table <- mof_table_df %>% knitr::kable(
  row.names = F,
  col.names = mof_table_col_names,
  format.args = list(scientific = F, digits = 4),
  booktabs = T,
  escape = F
)

# t-test----
# Standard error
se <- as.vector(sqrt(s2 * diag(solve(t(X) %*% X))))
# Vector of _t_ statistics
t_stats <- (b - H0) / se
# Calculate the p-values
if (two_tail == T) {
  p_values <- pt(q = abs(t_stats), df = dof, lower.tail = F) * 2
} else {
  p_values <- pt(q = abs(t_stats), df = dof, lower.tail = F)
}
# Do we (fail to) reject?
reject <- ifelse(p_values < alpha, reject <- "Reject", reject <- "Fail to Reject")

# Nice table (data.frame) of results
ttest_df <- data.frame(
  # The rows have the coef. names
  effect = rownames(b),
  # Estimated coefficients
  coef = as.vector(b) %>% round(3),
  # Standard errors

```

```

    std_error = as.vector(se) %>% round(3),
    # t statistics
    t_stat = as.vector(t_stats) %>% round(3),
    # p-values
    p_value = as.vector(p_values) %>% round(4),
    # reject null?
    significance = as.character(reject)
  )

  ttest_table <- ttest_df %>% knitr::kable(
    booktabs = T,
    format.args = list(scientific = F),
    escape = F
  )

  # Data frame for exporting for y, y_hat, X, and e vectors ----
  export_df <- data.frame(y, y_hat, e, X) %>% tbl_df()
  colnames(export_df) <- c("y", "y_hat", "e", colnames(X))

  # Return ----
  return(list(n=n, dof=dof, b=b, vars=export_df, R2uc=R2uc, R2=R2,
    R2adj=R2adj, AIC=AIC, SIC=SIC, s2=s2, SST=SST, SSR=SSR,
    mof_table=mof_table, ttest=ttest_table))
}

```

We'll also need a function to do an F-test for this Problem Set.

```

# Joint test function (from Ed's notes). Could also write a more complex functions that takes R and
F_test <- function(data, y_var, X_vars) {
  # Turn data into matrices
  y <- to_matrix(data, y_var)
  X <- to_matrix(data, X_vars)
  # Add intercept
  X <- cbind(1, X)
  # Name the new column "intercept"
  colnames(X) <- c("intercept", X_vars)
  # Calculate n and k for degrees of freedom
  n <- nrow(X)
  k <- ncol(X)
  # J is k-1
  J <- k - 1
  # Create the R matrix: bind a column of zeros
  # onto a J-by-J identity matrix
  R <- cbind(0, diag(J))

  # Estimate coefficients
  b <- b_ols(data, y_var, X_vars)
  # Retrieve OLS residuals
  e <- b$vars$e
  # Retrieve s^2

```

```

s2 <- b$s2 %<>% as.numeric()

# Create the inner matrix  $R(X'X)^{-1}R'$ 
RXXR <- R %*% solve(t(X) %*% X) %*% t(R)
# Calculate the F stat
f_stat <- t(R %*% b) %*% solve(RXXR) %*% (R %*% b) / J / s2
# Calculate the p-value;; why normal and not  $\chi^2$ 
p_value <- pf(q = f_stat, df1 = J, df2 = n-k, lower.tail = F)
# Create a data.frame of the f stat. and p-value
results <- data.frame(
  f_stat = f_stat %>% as.vector(),
  p_value = p_value %>% as.vector())
return(results)
}

```

Question 1:

Read the data into R. Inspect it. Sort by size (Q (kwh)).

```

nerlove <- readxl::read_excel("nerlove.xls", col_names=T)
summary(nerlove)

```

```

##          TC          Q          PL          PF
## Min.   : 0.082   Min.   :    2   Min.   : 1.450   Min.   :10.30
## 1st Qu.: 2.382   1st Qu.: 279   1st Qu.: 1.760   1st Qu.:21.30
## Median : 6.754   Median : 1109   Median : 2.040   Median :26.90
## Mean   : 12.976   Mean   : 2133   Mean   : 3.208   Mean   :26.18
## 3rd Qu.: 14.132   3rd Qu.: 2507   3rd Qu.: 2.190   3rd Qu.:32.20
## Max.   :139.422   Max.   :16719   Max.   :181.000   Max.   :42.80
##          PK
## Min.   :138.0
## 1st Qu.:162.0
## Median :170.0
## Mean   :174.5
## 3rd Qu.:183.0
## Max.   :233.0

```

```

# Fix typo in 13th row (missing a decimal!)
# DO THIS MORE ELEGANTLY!
nerlove[13, "PL"] <- 1.81

```

```

# nerlove %>%
#   filter(PL > 100) %>%
#   mutate(PL = PL/100)

```

```

nerlove %>% arrange(Q)

```

```

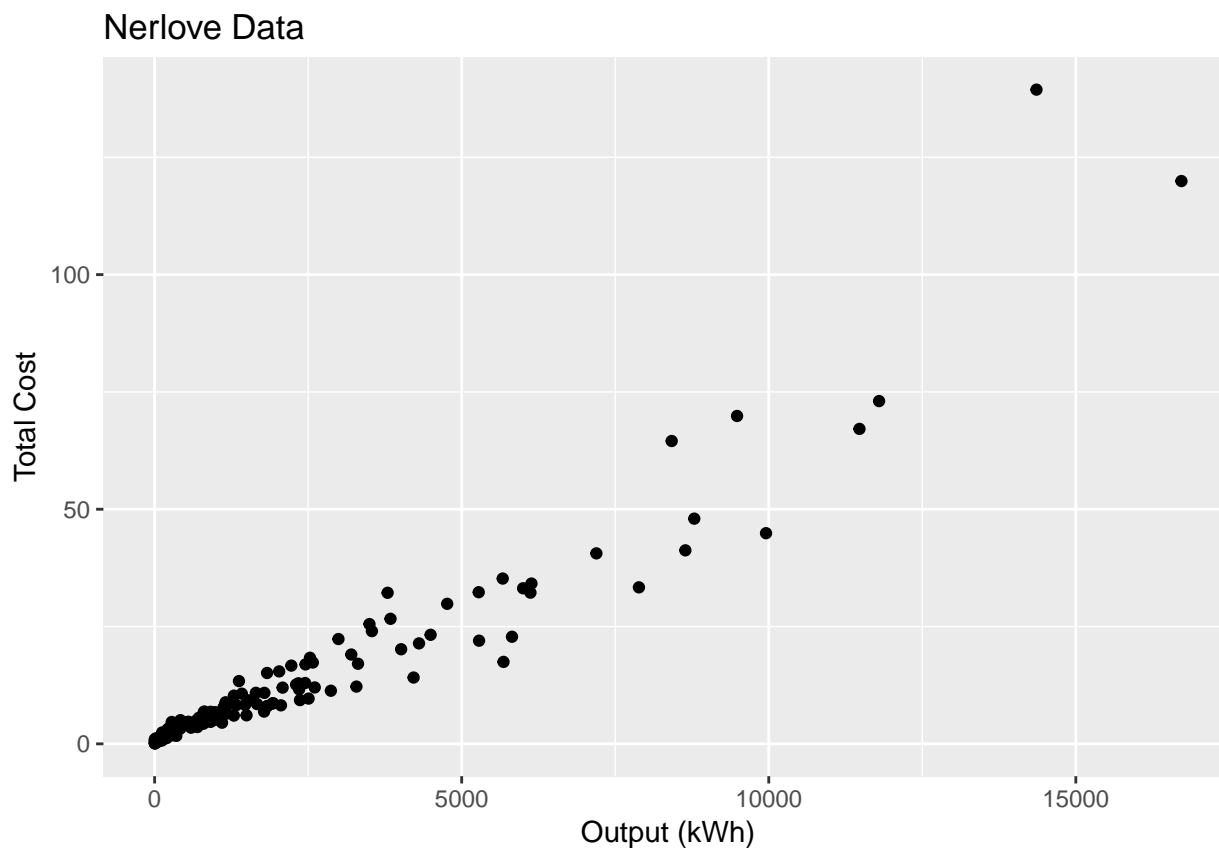
## # A tibble: 145 x 5

```

```
##      TC      Q      PL      PF      PK
##      <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.0820    2.   2.09  17.9  183.
## 2 0.661     3.   2.05  35.1  174.
## 3 0.990     4.   2.05  35.1  171.
## 4 0.315     4.   1.83  32.2  166.
## 5 0.197     5.   2.12  28.6  233.
## 6 0.0980    9.   2.12  28.6  195.
## 7 0.949    11.   1.98  35.5  206.
## 8 0.675    13.   2.05  35.1  150.
## 9 0.525    13.   2.19  29.1  155.
## 10 0.501   22.   1.72  15.0  188.
## # ... with 135 more rows
```

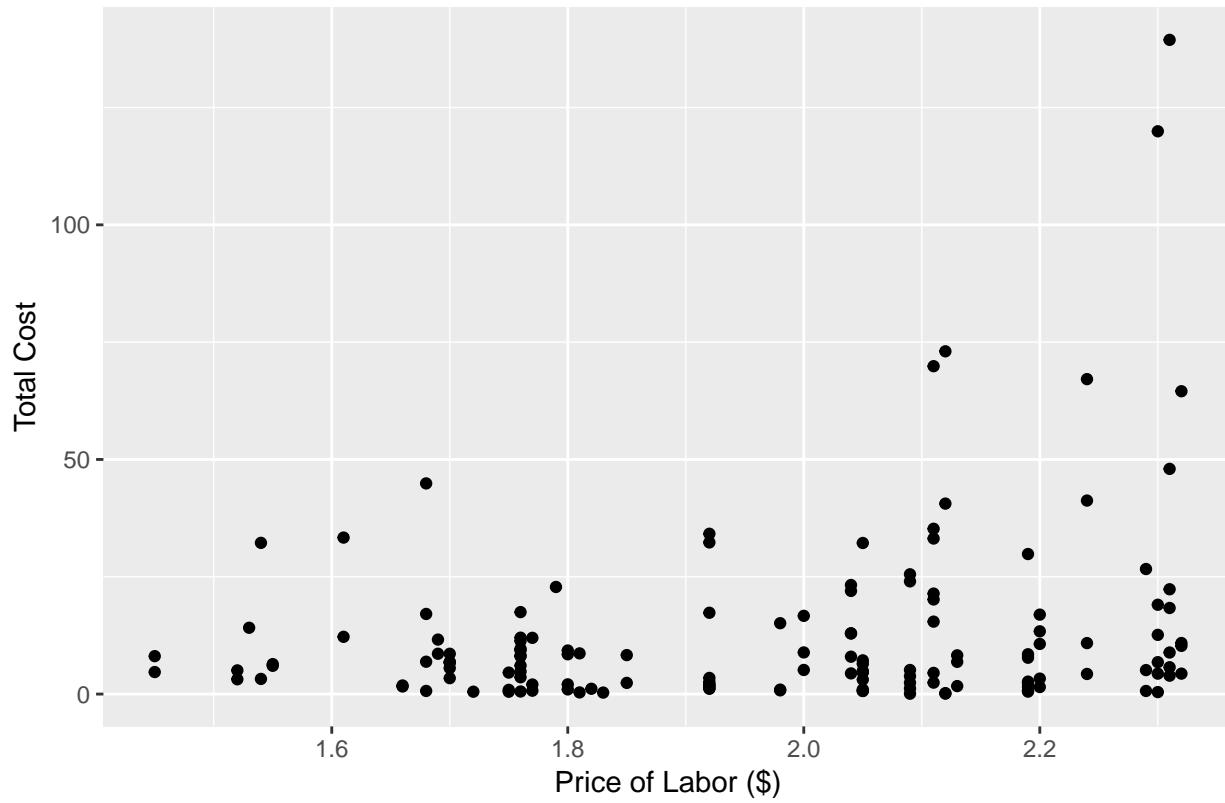
Plot the series and make sure your data are read in correctly.

```
ggplot(nerlove, aes(x=Q, y=TC)) +
  geom_point() +
  labs(title="Nerlove Data", x="Output (kWh)", y="Total Cost")
```



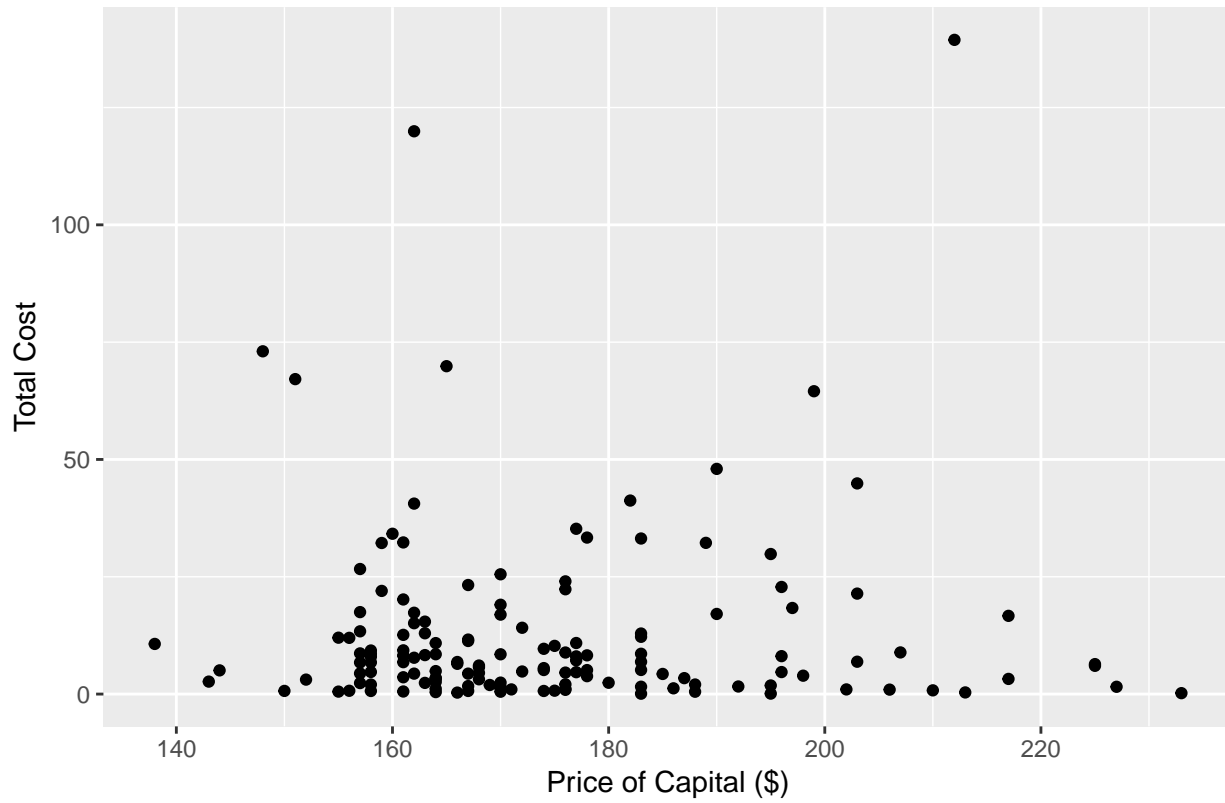
```
ggplot(nerlove, aes(x=PL, y=TC)) +
  geom_point() +
  labs(title="Nerlove Data", x="Price of Labor ($)", y="Total Cost")
```

Nerlove Data



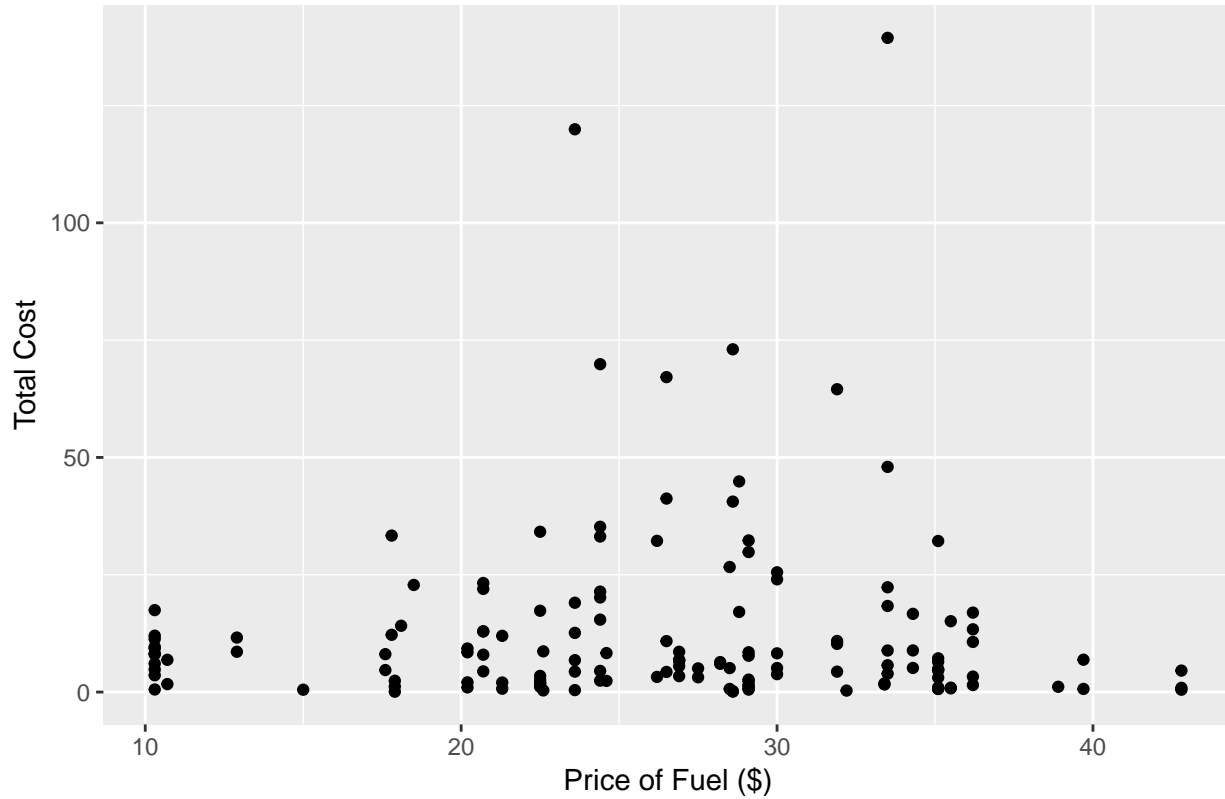
```
ggplot(nerlove, aes(x=PK, y=TC)) +  
  geom_point() +  
  labs(title="Nerlove Data", x="Price of Capital ($)", y="Total Cost")
```

Nerlove Data



```
ggplot(nerlove, aes(x=PF, y=TC)) +  
  geom_point() +  
  labs(title="Nerlove Data", x="Price of Fuel ($)", y="Total Cost")
```

Nerlove Data



Question 2:

Replicate regression I (page 176) in the paper.

Regression I:

$$\log(TC) - \log(P_F) = \beta_0 + \beta_1 Q + \beta_2 (\log(P_L) - \log(P_F)) + \beta_3 (\log(P_K) - \log(P_F))$$

Equivalent to:

$$\log\left(\frac{TC}{P_F}\right) = \beta_0 + \beta_1 Q + \beta_2 \log\left(\frac{P_L}{P_F}\right) + \beta_3 \log\left(\frac{P_K}{P_F}\right)$$

Where:

TC = total production cost,

P_L = wage rate,

P_K = “price” of capital,

P_F = price of fuel,

Q = output (measured in kWh)

In generalized Cobb-Douglas form:

$$\beta_1 = \frac{1}{r},$$

$$\beta_2 = \frac{a_L}{r},$$

$$\beta_3 = \frac{a_K}{r}$$

Prepare variables for Regression I.


```

# Create log variables
nerlove %<>% mutate(
  TClog = log(TC),
  Qlog = log(Q),
  PLlog = log(PL),
  PKlog = log(PK),
  PFlog = log(PF)
)

# Create PF scaled variables
nerlove %<>% mutate(
  TCscaled = TClog - PFlog,
  PLscaled = PLlog - PFlog,
  PKscaled = PKlog - PFlog
)

```

Variable names:

$$\log\left(\frac{TC}{P_F}\right) = \text{"TCscaled"}$$

$$\log\left(\frac{PL}{P_F}\right) = \text{"PLscaled"}$$

$$\log\left(\frac{PK}{P_F}\right) = \text{"PKscaled"}$$

```

# Regression I:
# dep var = (log costs - log fuel price) = TCscaled
reg_I <- ols(data = nerlove, y_data = "TCscaled",
             X_data = c("Qlog", "PLscaled", "PKscaled"),
             intercept = T, H0 = 0, alpha = 0.05)

reg_I$ttest

```

effect	coef	std_error	t_stat	p_value	significance
intercept	-2.037	0.384	-5.301	0.0000	Reject
Qlog	0.721	0.017	41.334	0.0000	Reject
PLscaled	0.593	0.205	2.898	0.0044	Reject
PKscaled	-0.007	0.191	-0.039	0.9692	Fail to Reject

```
reg_I$mof_table
```

R_{uc}^2	R^2	R_{adj}^2	SIC	AIC	SSR	s^2
0.966	0.9316	0.9301	-3.433	-3.515	4.082	0.02895

Coefficients are pretty close to those in the paper. R^2 matches!

Question 3:

Conduct the hypothesis test using constant returns to scale ($\beta_1 = 1$) as your null hypothesis.

```
# Re-run regression with null hypothesis H0 = 1
reg_I <- ols(data = nerlove, y_data = "TCscaled",
             X_data = c("Qlog", "PLscaled", "PKscaled"),
             intercept = T, H0 = 1, alpha = 0.05)
```

```
reg_I$ttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-2.037	0.384	-7.903	0.0000	Reject
Qlog	0.721	0.017	-16.020	0.0000	Reject
PLscaled	0.593	0.205	-1.990	0.0485	Reject
PKscaled	-0.007	0.191	-5.282	0.0000	Reject

What is the p- value associated with you test statistic? What is your point estimate of returns to scale? Constant? Increasing? Decreasing?


The p-value is 0.000. The point estimate of returns to scale is $\frac{1}{\beta} = r = 1.3875639$, hence returns to scale is increasing.

Question 4:

Plot residuals against output.

```
ggplot(reg_I$vars, aes(y=e, x=Qlog)) + geom_point() + labs(title="Regression I: Residuals against
```

Regression I: Residuals against Output



What do you notice? What does this potentially tell you from an economic perspective?

Evidence of heteroskedasticity: residuals seem to track the log output through parabola. We may want to rethink our specification!

Compute the correlation coefficient of the residuals with output for the entire sample? What does this tell you?

```
# R = cov(xy)/var(x)var(y)
R_I <- (cov(x=reg_I$vars$e, y=reg_I$vars$Qlog)/(var(reg_I$vars$e)*var(reg_I$vars$Qlog))) %>% signif(2)
R_I
```

```
## [1] "7.01e-14"
```

The correlation coefficient is extremely small: 7.01e-14.

Question 5:

Divide your sample into 5 subgroups of 29 firms each according to the level of output. Estimate the regression model again for each group separately.

```
# Divide sample into 5 subgroups
d <- split(nerlove,rep(1:5,each=29))

# Now we have a list, d, with five dataframes '1' through '5' for our five subgroups.

# Regression IIIA
reg_IIIA <- ols(data = d$`1`, y_data = "TCscaled",
               X_data = c("Qlog", "PLscaled", "PKscaled"),
               intercept = T, H0 = 1, alpha = 0.05)
reg_IIIA$ttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-1.452	1.366	-1.795	0.0848	Fail to Reject
Qlog	0.400	0.084	-7.101	0.0000	Reject
PLscaled	0.615	0.729	-0.528	0.6024	Fail to Reject
PKscaled	-0.081	0.706	-1.531	0.1384	Fail to Reject

```
# Regression IIIB
reg_IIIB <- ols(data = d$`2`, y_data = "TCscaled",
               X_data = c("Qlog", "PLscaled", "PKscaled"),
               intercept = T, H0 = 1, alpha = 0.05)
reg_IIIB$ttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-2.818	0.614	-6.222	0.0000	Reject
Qlog	0.658	0.116	-2.939	0.0070	Reject
PLscaled	0.094	0.274	-3.304	0.0029	Reject
PKscaled	0.378	0.277	-2.250	0.0335	Reject

Regression IIIC

```
reg_IIIC <- ols(data = d$`3`, y_data = "TCscaled",
               X_data = c("Qlog", "PLscaled", "PKscaled"),
               intercept = T, H0 = 1, alpha = 0.05)
reg_IIIC$ttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-3.185	0.734	-5.705	0.0000	Reject
Qlog	0.938	0.198	-0.312	0.7578	Fail to Reject
PLscaled	0.402	0.199	-2.997	0.0061	Reject
PKscaled	0.250	0.187	-4.010	0.0005	Reject

Regression IIID

```
reg_IIID <- ols(data = d$`4`, y_data = "TCscaled",
               X_data = c("Qlog", "PLscaled", "PKscaled"),
               intercept = T, H0 = 1, alpha = 0.05)
reg_IIID$ttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-2.843	0.506	-7.596	0.0000	Reject
Qlog	0.912	0.107	-0.818	0.4210	Fail to Reject
PLscaled	0.507	0.187	-2.630	0.0144	Reject
PKscaled	0.093	0.164	-5.525	0.0000	Reject

Regression IIIE

```
reg_IIIE <- ols(data = d$`5`, y_data = "TCscaled",
               X_data = c("Qlog", "PLscaled", "PKscaled"),
               intercept = T, H0 = 1, alpha = 0.05)
reg_IIIE$ttest
```

effect	coef	std_error	t_stat	p_value	significance
intercept	-2.916	0.454	-8.618	0.0000	Reject
Qlog	1.044	0.065	0.683	0.5008	Fail to Reject
PLscaled	0.603	0.197	-2.014	0.0549	Fail to Reject
PKscaled	-0.289	0.175	-7.374	0.0000	Reject

May want to clean this up and run as for loop???

Can you replicate Equations IIIA - IIIE? Calculate the point estimates for returns to scale for each sample. Is there a pattern relating to size of output? Regression | Returns to Scale
 ———— | ————— IIIA | 2.4981863 IIIB | 1.5194085 IIIC | 1.0657807 IIID | 1.0964381 IIIE | 0.9574966

Coefficients roughly match the results of the paper!

Question 6:

Create "dummy variables" for each industry. Interact them with the output variable to create five "slope coefficients".

```
# create group categorical var
for (i in 1:5) {
  d[[i]] %<>% mutate(gvar=i)
}

# unsplit
df <- rbind(d[[1]], d[[2]], d[[3]], d[[4]], d[[5]])

# create dummies
df %<>%
  mutate(
    g1 = ifelse(gvar==1, 1, 0),
    g2 = ifelse(gvar==2, 1, 0),
    g3 = ifelse(gvar==3, 1, 0),
    g4 = ifelse(gvar==4, 1, 0),
    g5 = ifelse(gvar==5, 1, 0))

df %<>% select(-gvar)

# interact with output variable
df %<>%
  mutate(
    lQ_1 = Qlog*g1,
    lQ_2 = Qlog*g2,
    lQ_3 = Qlog*g3,
    lQ_4 = Qlog*g4,
    lQ_5 = Qlog*g5)
```

Run a model, letting the intercept and slope coefficient on output differ across plants, but let the remainder of the coefficients be pooled across plants.

I'm not sure this is the right model... but I'm gettin similar coeffs!

```
reg_IV <- ols(data = df, y_data = "TCscaled",
  X_data = c("lQ_1", "lQ_2", "lQ_3", "lQ_4", "lQ_5", "g1",
    "g2", "g3", "g4", "g5", "PLscaled", "PKscaled"),
  intercept = F, H0 = 1, alpha = 0.05)

# Can also omit one of the dummie vars (g) and include an intercept >> get the same coeffs
reg_IV$ttest
```

effect	coef	std_error	t_stat	p_value	significance
lQ_1	0.397	0.043	-14.002	0.0000	Reject
lQ_2	0.648	0.147	-2.389	0.0183	Reject

effect	coef	std_error	t_stat	p_value	significance
lQ_3	0.885	0.297	-0.388	0.6989	Fail to Reject
lQ_4	0.909	0.274	-0.333	0.7393	Fail to Reject
lQ_5	1.063	0.131	0.478	0.6336	Fail to Reject
g1	-1.815	0.305	-9.231	0.0000	Reject
g2	-2.194	0.489	-6.538	0.0000	Reject
g3	-2.879	0.972	-3.993	0.0001	Reject
g4	-2.922	0.966	-4.059	0.0001	Reject
g5	-3.511	0.599	-7.525	0.0000	Reject
PLscaled	0.426	0.163	-3.520	0.0006	Reject
PKscaled	0.104	0.152	-5.888	0.0000	Reject
Note: Numbers 1 through 5 correspond to A through E					

Are there any noticeable changes in returns to scale from the previous part?

yes. . . .

Question 7:

Conduct a statistical test comparing the first model you estimate to the last model you estimated. (Hint: Is one model a restricted version of the other?). Would separate t-test have given you the same results?

Question 8:

To see whether returns to scale declined with output, Nerlove tested a nonlinear specification by including $\ln(y)^2$ as a regressor. Conduct a statistical test you feel is appropriate to test this hypothesis.